

UNIT 2

Messaging & Positioning

Crafting Messages for Humans and AI

Future of PMM Academy

futureofproductmarketing.com

Unit 2: Messaging & Positioning in the AI Era

The Two-Audience Problem

Every piece of content you create now serves two audiences: humans and machines. This isn't a future state — it's the current reality, and most PMM teams are still writing exclusively for the first audience while ignoring the second.

When you craft a positioning statement, a human reader processes it through emotional resonance, brand affinity, and personal relevance. They might not remember your exact words, but they'll remember how you made them feel. Traditional messaging guidance optimizes for this: be memorable, be differentiated, create an emotional connection.

But an AI reading that same statement is doing something different. It's evaluating semantic content. It's cross-referencing claims against other sources. It's assessing specificity, looking for verifiable details, and deciding whether to cite you in a synthesized response or pass over you for a competitor who made a more concrete claim.

These two optimization targets don't always conflict — but they often do. A tagline like "Empowering Digital Transformation" might resonate emotionally with a human executive who's tired of legacy systems. But it gives an AI nothing to work with. No specificity. No verifiable claim. No differentiation from the thousand other vendors who say the same thing.

The 10x PMM solves the two-audience problem by creating messaging that works for both: emotionally resonant for humans, informationally dense for machines.

Section 1: The Specificity Imperative

Vague Claims Vanish

The single biggest shift in messaging strategy for the AI era is the move from vague to specific.

Compare these two value propositions:

Vague: "Our platform helps teams work faster and smarter."

Specific: "Teams using our platform complete quarterly close 40% faster, with audit findings reduced by 65% in the first year."

The vague claim is technically true of every software product ever built. It provides no differentiation. When an AI agent is asked to compare accounting platforms, it has nothing to cite.

The specific claim gives the AI something to surface. It includes numbers. It implies verifiable outcomes. It can be cross-referenced. And crucially, it can be used in a comparison: "Platform A cites 40% faster quarterly close; Platform B does not provide comparable metrics."

This doesn't mean every sentence needs statistics. But it does mean your core positioning — the claims you want buyers to remember and AI agents to surface — must be specific enough to stand out from the noise of generic B2B messaging.

The Specificity Ladder

Think of specificity as a ladder with five rungs:

1. **Abstract:** "We make things better." (Zero specificity — invisible to AI)
2. **Categorical:** "We improve productivity." (Category-level — AI surfaces but doesn't differentiate)
3. **Directional:** "We improve productivity for sales teams." (Audience-specific — AI can filter)
4. **Quantitative:** "We improve sales productivity by 30%." (Measurable — AI can compare)

5. **Evidenced:** “We improve sales productivity by 30%, based on a 2025 study of 200 enterprise deployments.” (Verifiable — AI trusts and cites)

Most B2B messaging lives at level 2 or 3. The companies winning the AI visibility game operate at level 4 and 5.

This requires a cultural shift in how PMMs work with stakeholders. When a product manager says “our new feature improves performance,” the PMM’s job is to push: “By how much? In what conditions? Can we measure it? Can we publish it?” That negotiation becomes a core PMM competency.

Section 2: Claims That Cite

The Citation Hierarchy

When an AI constructs a response about your category, it’s implicitly ranking sources. Understanding that hierarchy helps you position your content to get cited.

Tier 1: Independent evaluation. Analyst reports (Gartner, Forrester, IDC), peer review sites (G2, TrustRadius), and academic or journalistic sources. AI models weight these heavily because they appear independent of vendor bias.

Tier 2: Documented facts. Your published specifications, pricing (if public), documentation, and any claims that are verifiable against third-party sources.

Tier 3: Company claims. Your website copy, press releases, and marketing materials. AI models treat these as potentially biased but still cite them when they provide specific information not available elsewhere.

Tier 4: Promotional content. Vague value propositions, emotional branding, lifestyle imagery. AI models tend to ignore this content entirely.

The practical implication: your goal is to move as much of your positioning as possible into Tiers 1 and 2. That means publishing data, supporting analyst research, encouraging customer reviews, and documenting everything.

The Documentation Layer

The most underrated asset in AI-era messaging is your documentation corpus.

Product documentation — help articles, technical specs, API references, implementation guides — is often treated as a post-sale concern, owned by technical writers rather than product marketers. This is a mistake. In the agentic buyer journey, documentation is pre-sale content.

When a buyer asks an AI to compare integration capabilities across vendors, the AI reads your documentation. All of it. It surfaces what’s specific and verifiable. If your docs are comprehensive and well-structured, you get cited. If they’re sparse or confusing, you get skipped.

This means PMM needs a seat at the documentation table. Not to rewrite technical content with marketing spin — that would backfire — but to ensure the strategic claims you’re making in positioning are substantiated somewhere in the verifiable corpus.

Section 3: Positioning for Discovery

Category Design in the AI Era

Category design — the strategic practice of defining and owning a market category — gets harder and easier simultaneously in the AI era.

Harder, because AI models are trained on existing category definitions. If you’re trying to create a new category, the AI doesn’t know it exists yet. A buyer who asks “what are the best CDP alternatives” will get answers about existing CDP vendors, not about your new category that you’ve decided makes CDPs obsolete.

Easier, because specificity wins. If your new category is genuinely different, and you can describe that difference with specific, verifiable claims, AI will eventually surface it. The model will learn that when someone asks about customer data,

there's this other thing that works differently. Your job is to provide enough specific content that the category becomes real in the AI's world model.

The tactical advice: don't abandon category design, but supplement it with relentless specificity. Define your category, but also define exactly how you differ from adjacent categories in measurable ways. Don't just say "we're not a CRM" — say "unlike CRMs that require manual data entry, our platform automatically ingests 47 data sources with no schema configuration."

Building a Citation Moat

A citation moat is the competitive advantage you build by being the most citable source in your category.

It works like this: the more often AI agents cite you in responses about your category, the more your claims get reinforced in buyer research, the more third parties reference those claims, the more training data exists with your positioning embedded in it. It's a flywheel.

To build a citation moat:

1. **Publish original data.** Industry benchmarks, survey results, technical comparisons. This is Tier 1 content that AI models treat as authoritative.
 2. **Be exhaustively specific.** For every feature you claim, document exactly what it does, how it works, and what outcomes it produces. Leave nothing implied.
 3. **Enable third-party validation.** Make it easy for customers to review you. Support analyst research with data. Provide quotes and statistics that journalists and bloggers can cite.
 4. **Update constantly.** AI models learn continuously. A claim you made in 2024 may be overwritten by fresher competitor content in 2026. Your messaging is never finished.
 5. **Structure for parsing.** Use comparison tables, spec sheets, and structured data wherever possible. AI models can process structured information more reliably than prose.
-

Section 4: Testing Messaging Across Audiences

The AI Message Test

Before launching any major positioning, test it against AI.

The test is simple: ask multiple AI models to compare you to competitors, explain your differentiators, or recommend solutions for buyer problems you're trying to solve. Note what gets surfaced, what gets missed, and what gets wrong.

This isn't a replacement for human validation — you still need to test messaging with actual buyers, sales teams, and analysts. But it's an additional validation layer that reveals whether your positioning is machine-readable, not just human-readable.

A structured testing protocol:

1. **Baseline.** Ask AI to describe your product and competitors with no prompting. What does it know today?
2. **Comparison.** Ask AI to compare your product to specific competitors. What differentiators does it surface?
3. **Recommendation.** Ask AI what solution it would recommend for your target buyer's problem. Are you on the list?
4. **Objection.** Ask AI about weaknesses or concerns with your product. What are buyers hearing before they talk to you?
5. **Source check.** Ask AI where it's getting its information. What sources is it citing?

Run this protocol quarterly, or after any major messaging change. Treat it as a leading indicator of how the agentic buyer journey will present your company.

Balancing Human and Machine

The goal is not to write only for machines. Cold, robotic messaging full of specs and statistics will fail with human buyers even if AI loves it.

The balance:

Headlines and taglines — optimize for human emotion and memorability. This is where you earn attention.

Body copy and explanations — blend emotion with specificity. Every emotional claim should be backed by specific evidence within a few sentences.

Documentation and specs — optimize for machine parsing. Structure, detail, verifiability.

Customer proof points — optimize for citation. Named customers, quoted metrics, third-party validation.

Think of it as layers. The human reads the surface layer and feels resonance. The AI reads the deeper layers and finds substance. Both leave with the impression you want them to have.

Key Takeaways

- **Every piece of content now serves two audiences** — humans who respond to emotion and AI agents who evaluate specificity.
 - **Vague claims vanish.** The specificity imperative means moving from “we make things better” to measurable, verifiable outcomes.
 - **Citations are the new conversions.** Tier 1 (independent) and Tier 2 (documented) claims get surfaced; Tier 4 (promotional fluff) gets ignored.
 - **Documentation is pre-sale content.** Your technical corpus is being read by AI before any human buyer sees your website.
 - **Build a citation moat.** Original data, exhaustive specificity, third-party validation, constant updates, and structured formats create compounding advantage.
 - **Test messaging against AI.** Run baseline, comparison, recommendation, objection, and source checks before launching positioning.
-

Hands-On Exercise

Apply the specificity ladder to your current positioning:

1. Write down your three core value propositions exactly as they appear today.
2. Score each one on the specificity ladder (1 = Abstract, 5 = Evidenced).
3. For any proposition scoring below 4, draft a more specific version.
4. Test both versions with an AI — ask it to compare you to a competitor using only public information. Which version gets surfaced?

Then, run a competitive citation check:

1. Ask an AI: “What are the key differences between [your product] and [top competitor]?”
 2. Note which of your differentiators it surfaces.
 3. For any missing differentiators, identify where in your public corpus that claim should be documented.
 4. Create a task to add or improve that documentation.
-

Next: Unit 3 — Competitive Intelligence & AI

